

The mystery of Time in Distributed Systems

Part - III


Cassandra don't need vector clocks

used for conflict resolution
How???

reference - data stax blog


Key = mandeep

values $v_0 = \{ \text{'email': 'msdeep14.ms@gmail.com'}, \text{'phone': 'nhi bataunga'} \}$



↓ update email address

$v_1 = \{ \text{'email': 'msdeep14.ms@msdeepSingh.com'}, \text{'phone': 'nhi bataunga'} \}$



↓ update phone no.

$v_2 = \{ \text{'email': 'msdeep14.ms@gmail.com'}, \text{'phone': '79825n3xyy'} \}$

read from replicat

and update it.

In general, last write wins algorithm picks most recent

write → V1 is discarded and V2 is picked up.
data loss

For vector clock, both V1 and V2 are retained.

↓
return both versions to client on next read and client can merge as appropriate.

Issues with Vector Clocks

- ① Performance → updating a field requires
 - ↳ read & deserialize existing object
 - ↳ update field
 - ↳ serialize and write back
- ② Siblings → multiple versions generated by conflicting updates ← difficult to deal
- ③ vector clock don't actually specify how to resolve the conflicts.

Solution in place @ Cassandra

↳ rows is broken into columns and updated independently.

```
CREATE TABLE users (  
  username text PRIMARY KEY,  
  email text,  
  phone text  
);  
  
INSERT INTO users (username, email, phone)  
VALUES ('jbellis', 'jbellis@example.com', '555-5555');  
  
UPDATE users SET email = 'jbellis@illustration.com' WHERE username = 'jbellis';  
  
UPDATE users SET phone = '444-4444' WHERE username = 'jbellis';
```

← image from blog

Benefits

① conflicts can be resolved automatically.

basis timestamp



Each column has its own

timestamp value.

See you in next video

Subscribe for more

You Tube - MsDeep Singh

Happy Learning 😊